

Advanced Java: Language, Internals, Techniques

Bob Singh

Learning JNI

Texts:

- **JNI**
 - **Java Native Interface Sheng Liang**
 - * java.sun.com/docs/books/jni/
 - **JDC Training Writing Advanced Applications**
 - * Chap 5 JNI Examples
 - **JNI FAQ**
 - * java.sun.com/products/jdk/faq/jnifaq.html

Multilanguage Java Programming

Motivation

- **Wrappers to legacy platform libraries e.g Java wrappers for MPI, OpenGL etc**
- **Embedding JVM in Server environments for additional functionality**
- **Performance when writing time-critical portion in low-level efficient languages**

Multilanguage Java Programming

Issues

- **Type-safety of Java can be compromised**
 - Extra care needed
- **Portability of Java not there**
 - May not run on multiple host environments

JNI

Objectives

- **Compatibility**
- **Compromise between VM independence and efficiency**
- **Functionality**
 - **Expose sufficient JVM details**

Learning JNI

JNI functionality

- **Calling into External Native Libraries**
 - Legacy Libraries
- **Embedding JVM into application**
 - Web Browsers, Web Servers, Databases etc
- **Some JVM exposure**
 - GetCritical etc stops GC
 - Does VM create a copy of Object or not?

Learning JNI

JNI Overview

- **From Java code**
 - Call native code written in C, C++
- **From C/C++ code**
 - Create, Update Java objects
 - Call Java Methods
 - Load classes and obtain class information

Learning JNI

Data types

- **Primitive types map straightforward**
 - Java type `int` maps to JNI C/C++ type `jint`
- **Reference types passed as opaque references**
 - C pointer types that refer to internal data structures in the JVM
 - Manipulate them through JNI functions

Learning JNI

HelloWorld

```
class HelloWorld {  
    private native void print();  
    public static void main(String[] args) {  
        new HelloWorld().print();  
    }  
    static {  
        System.loadLibrary("HelloWorld");  
    }  
}
```

Learning JNI

HelloWorld

```
#include <jni.h>
#include <stdio.h>
#include "HelloWorld.h"

JNIEXPORT void JNICALL
Java_HelloWorld_print(JNIEnv *env, jobject obj)
{
    printf("Hello World!\n");
    return;
}
```

Learning JNI

HelloWorld

Solaris

```
cc -G -I/java/include -I/java/include/solaris  
HelloWorld.c -o libHelloWorld.so
```

Linux

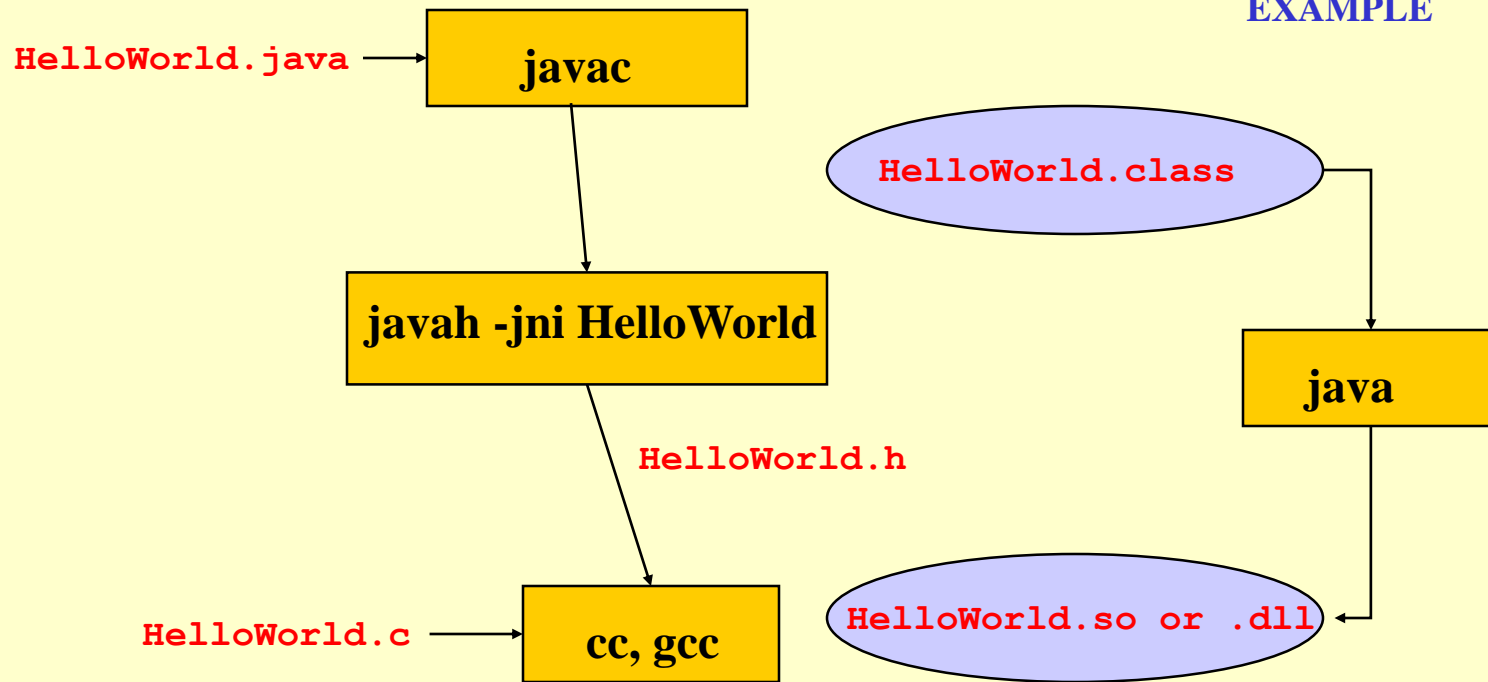
```
gcc -o libHelloWorld.so -shared -Wl,-  
soname,libHelloWorld.so -I/home/jdk1.2/include -  
I/home/jdk1.2/include/linux HelloWorld.c -static -lc
```

Win32

```
cl -Ic:\java\include -Ic:\java\include\win32 -MD -LD  
HelloWorld.c -FeHelloWorld.dll
```

Learning JNI

EXAMPLE



Learning JNI

Invocation Overview

```
JavaVM *jvm;  
JNIEnv *jni_env;  
  
JNI_CreateJavaVM(&jvm, &jni_env, &args);  
  
jclass cls = jni_env->FindClass("Hello");  
jmethodID mid = jni_env->GetStaticMethodID(cls, "run",  
"()V");  
jni_env->CallStaticVoidMethod(cls, mid);  
  
jvm->DestroyJavaVM();
```

Learning JNI

Data types

- **jobject** supertype
- subtypes **jstring**, **jclass**, **jarray**
 - Commonly used reference types
 - **jstring** corresponds to **java.lang.String**
 - **jclass** corresponds to **java.lang.Class**
 - **jobjectarray**, **jbooleanarray** etc are subtypes of **jarray**

Learning JNI

String types

- **java strings Unicode**
- **C strings in byte format**
- **UTF allows encoding of Unicode (Unicode Transformation Format)**
 - **non-null ASCII represented as C byte (8 bits)**
 - * **\u0001 to \u007f represented by 8 bits**
 - * **\u0080 to \u07ff represented by 16 bits**
 - * **\u0800 to \uffff represented by 24 bits**

Learning JNI

EXAMPLE

`\u0001 to \u007f` →

0	
---	--

`\u0080 to \u07ff` →

1	1	0	
---	---	---	--

1	0	
---	---	--

`\u0800 to \uffff` →

1	1	1	0	
---	---	---	---	--

1	0	
---	---	--

1	0	
---	---	--

Learning JNI

Other String Functions

- **JDK 1.1.x and Java2**
 - GetStringChars, ReleaseStringChars, GetStringLength
 - GetStringUTFChars, ReleaseStringUTFChars, GetStringUTFLength
 - NewString, NewStringUTF
- **Java2**
 - GetStringCritical, ReleaseStringCritical
 - GetStringRegion, GetStringUTFRegion

Learning JNI

Prompt

```
class Prompt {  
  
    // native method that prints a prompt and reads a line  
    private native String getLine(String prompt);  
  
    public static void main(String args[]) {  
        Prompt p = new Prompt();  
        String input = p.getLine("Type a line: ");  
        System.out.println("User typed: " + input);  
    }  
  
    static {  
        System.loadLibrary("Prompt");  
    }  
}
```

Learning JNI

Prompt

```
#include <jni.h>
#include <stdio.h>
#include "Prompt.h"

JNIEXPORT jstring JNICALL
Java_Prompt_getLine(JNIEnv *env, jobject obj, jstring prompt)
{
    char buf[128];
    const char *str;
    str = (*env)->GetStringUTFChars(env, prompt, NULL);
    if (str == NULL) {
        return NULL; /* OutOfMemoryError already thrown */
    }
    printf("%s", str);
    (*env)->ReleaseStringUTFChars(env, prompt, str);
    /* We assume here that the user does not type more than
       * 127 characters */
    scanf("%s", buf);
    return (*env)->NewStringUTF(env, buf);
}
```

Learning JNI

Prompt

```
#include <jni.h>
#include <stdio.h>
#include "Prompt.h"

JNIEXPORT jstring JNICALL
Java_Prompt_getLine(JNIEnv *env, jobject obj, jstring prompt)
{
    char buf[128];
    const char *str;
    str= prompt;
    printf("%s", str);
    scanf("%s", buf);
    return (*env)->NewStringUTF(env, buf);
}
```

Learning JNI

Prompt

C

```
str = (*env)->GetStringUTFChars(env, prompt, NULL);
```

C++

```
str = env->GetStringUTFChars(env, prompt, NULL);
```

Learning JNI

Array types

- **Primitive Arrays**

- **Similar to Strings**

- **JDK 1.1.x**

- * **Get<Type>ArrayRegion, Set<Type>ArrayRegion**

- * **Get<Type>ArrayElements, Release<Type>ArrayElements**

- * **GetArrayLength**

- * **New<Type>Array**

- **Java2 (finer control)**

- * **GetPrimitiveArrayCritical, ReleasePrimitiveArrayCritical**

- **Object Arrays**

- * **GetObjectArrayElement, SetObjectArrayElement**

Learning JNI

Summation

```
class Summation {  
  
    // native method that sums the given array  
    private native int sum(int[] arr);  
  
    public static void main(String args[]) {  
        Summation p = new Summation();  
        int arr[] = new int[10];  
        for(int j=0; j < arr.length; j++)  
            arr[j] = j;  
        int sum = p.sum(arr);  
        System.out.println("Sum = " + sum);  
    }  
  
    static {  
        System.loadLibrary("Summation");  
    }  
}
```

Learning JNI

Summation Version #1

```
#include <jni.h>
#include <stdio.h>
#include "Summation.h"

JNIEXPORT jint JNICALL
Java_Summation_sum(JNIEnv *env, jobject obj, jintArray arr)
{
    jint buf[10];
    jint j, sum = 0;
    (*env)->GetIntArrayRegion(env, arr, 0, 10, buf);
    for(j=0; j < 10; j++)
        sum += buf[j];
    return sum;
}
```


Learning JNI

Summation Version #2

```
#include <jni.h>
#include <stdio.h>
#include "Summation.h"

JNIEXPORT jint JNICALL
Java_Summation_sum(JNIEnv *env, jobject obj, jintArray arr)
{
    jint *carr;
    jint j, sum = 0;
    carr = (*env)->GetIntArrayElements(env, arr, NULL);
    if (carr == NULL)
        return 0;
    for(j=0; j < 10; j++)
        sum += carr[j];
    (*env)->ReleaseIntArrayElements(env, arr, carr, 0);
    return sum;
}
```

Learning JNI

Field Access

```
class StringArrayInit {  
  
    // native method that initializes an Array of Strings  
    private native String[] initStringArr();  
  
    public static void main(String args[]) {  
        StringArrayInit p = new StringArrayInit();  
        String[] arr = p.initStringArr();  
        for(int j=0; j < arr.length; j++) {  
            System.out.println(arr[j]);  
        }  
    }  
  
    static {  
        System.loadLibrary("StringArrayInit");  
    }  
}
```

Learning JNI

ObjectInitialization

```
#include <jni.h>
#include <stdio.h>
#include "StringArrayInit.h"

JNIEXPORT jobjectArray JNICALL
Java_StringArrayInit_initStringArr(JNIEnv *env, jobject obj) {
    jobjectArray ret;
    jint j;
    char *msg[3] = {"abc", "def", "ghi"};
    jclass strCls = (*env)->FindClass(env, "java/lang/String");
    jstring str = (*env)->NewStringUTF("");
    ret = (*env)->NewObjectArray(3, strCls, str);
    for(j=0; j < 3; j++)
        (*env)->SetObjectArrayElement(ret, j,
                                        (*env)->NewStringUTF(msg[j]));
    return ret;
}
```

Learning JNI

Fields and Methods

- **Fields**

- **Instance**

- * **GetFieldID**

- * **Get<Type>Field** and **Set<Type>Field**

- **Static**

- * **GetStaticFieldID**

- * **GetStatic<Type>Field** and **SetStatic<Type>**

- **Methods**

- * **Instance: GetMethodID, Call<Type>Method**

- * **Static: GetStaticMethodID, CallStatic<Type>Method**

Learning JNI

Method Calling Pattern

- **Retrieve the class**
 - FindClass through name
 - GetObjectClass from a jobject
 - Directly from a jclass
- **Retrieve the MethodID**
 - GetMethodID, or GetStaticMethodID
- **Call the Method**
 - Using Call<Type>Method

Learning JNI

Method Calling

- **Constructors**
 - “<init>” method-name and “V” return type
- **Method Descriptor**
 - Use `javap -s`
- **SuperClass methods**
 - `CallNonVirtual<Type>Method`

Learning JNI

FieldAccess

```
class InstanceFieldAccess {
    private String s;

    private native void accessField();
    public static void main(String args[]) {
        InstanceFieldAccess c = new InstanceFieldAccess();
        c.s = "abc";
        c.accessField();
        System.out.println("In Java:");
        System.out.println("  c.s = \"" + c.s + "\"");
    }
    static {
        System.loadLibrary("InstanceFieldAccess");
    }
}
```

Learning JNI

FieldAccess

```
#include <jni.h>
#include <stdio.h>
#include "InstanceFieldAccess.h"

JNIEXPORT void JNICALL
Java_InstanceFieldAccess_accessField(JNIEnv *env, jobject obj)
{
    jfieldID fid;    /* store the field ID */
    jstring jstr;
    const char *str;
    jclass cls = (*env)->GetObjectClass(env, obj);
    printf("In C:\n");
    /* Look for the instance field s in cls */
    fid = (*env)->GetFieldID(env, cls, "s",
                             "Ljava/lang/String;");

    /* Read the instance field s */
    jstr = (*env)->GetObjectField(env, obj, fid);
    str = (*env)->GetStringUTFChars(env, jstr, 0);
    printf("  c.s = \"%s\"\n", str);
    (*env)->ReleaseStringUTFChars(env, jstr, str);
    jstr = (*env)->NewStringUTF(env, "123");
    (*env)->SetObjectField(env, obj, fid, jstr);
}
```


Learning JNI

Method and Field Lookup

- **Expensive**
- **Use Caching for Field and Method ID's**
 - In the static initializer setup initID's
- **Need to know about other things for caching**
 - Stale ID's if Class gets Unloaded and Reloaded

Learning JNI

Opaque References

- **C pointer types that refer to JVM internal data structures, manipulated through JNI functions**
 - **Local**
 - * Valid for duration of native method
 - * Illegal to use across threads
 - * Illegal to cache across invocations
 - * `NewObject`, `FindClass` etc return local Refs
 - **Global**
 - * Valid until programmer frees
 - **Weak Global**
 - * new in Java2

Learning JNI

GlobalRefs from java.net

```
static jclass socketExceptionCls;
JNIEXPORT void JNICALL
Java_java_net_PlainSocketImpl_socketCreate(JNIEnv *env, jobject o) {

    if (socketExceptionClass == NULL) {
        socketExceptionCls = (*env)->FindClass(env,
"java/net/SocketException");
        socketExceptionCls = (jclass)(*env)->NewGlobalRef(env,
socketExceptionCls);
        :
    }
}
```

Learning JNI

LocalRef Handling in Java2

```
#define NREFS

for(j=0; j < len; j++) {
    if ((*env)->PushLocalFrame(env, NREFS) < 0) {
        Error
    }
    jstr = (*env)->GetObjectArrayElement(env, arr, j);
    /* Process jstr */

    (*env)->PopLocalFrame(env, NULL);
}
```

Learning JNI

Exceptions

- **Exception condition encoded in JNI API**
- **ExceptionCheck**
 - Needed for **Call<Type>Method**
- **ExceptionDescribe, ExceptionClear**
- **ThrowNew**
 - Throwing New Exceptions

Learning JNI

JNI Gotchas

- **Expensive on some JVM's**
- **Need extensive Error Checking**
- **JNIEnv, Local Refs etc are per thread**
 - **Don't use across threads by caching**
- **Threading Models should match**

Learning JNI

JNI Java2 tips

- **verbose:jni option**
- **-Xcheck:jni**
 - Very expensive
- **EnsureLocalCapacity, PushLocalFrame, PopLocalFrame**
- **GetStringCritical, GetPrimitiveArrayCritical**
 - Direct pointer to JVM possible

JNI Performance Issues

Calls to Native Methods

- **Function/method calls in general are costly**
 - Creation of new frame on execution stack
 - Save/restore of current function's state
- **Conversion of *calling convention***
 - Method implementation may not be known at compile time
 - Interpreter and JIT calling conventions traditionally designed for fast method calling within their respective worlds
 - Newer systems must optimize paths between these worlds

JNI Performance Issues

Accessing Arrays in Native Code

- **Directly in Java object space**
 - Garbage collector cannot move object until native code has finished
 - * Object can be pinned
 - * GC can be disabled
- **Copying array into native memory space**
 - Object must be copied twice to access and modify

JNI Performance Issues

JNI Array Access Routines

- **Get<Type>ArrayRegion**
 - Obtain pointer to copy of specified region of array
 - Any changes can be committed to Java object with Set<Type>ArrayRegion call if desired
- **Get<Type>ArrayElements**
 - Gets entire array, changes committed to Java object at accompanying Release<Type>ArrayElements call
 - Allows VM to choose implementation
 - * Pinning: most Sun-based VM's
 - * Copying: Sun Production VM, HotSpot

JNI Performance Issues

JNI Array Access Routines

- **GetPrimitiveArrayCritical**
 - Gets entire array, changes committed to Java object at accompanying **SetPrimitiveArrayCritical** call
 - Instructs VM to avoid copying if at all possible
 - * Pinning
 - * Disabling garbage collector
 - **Restrictions while array is held**
 - * No JNI calls
 - * No blocking operations


JNI Benchmarking

Goals

- **Measurement of all important native functionality**
 - Calls into native code, parameter passing
 - Native code access to Java fields and methods
 - Native code access to Java arrays
 - Native code management of Java references (global, local, weak, ...)
- ***Microbenchmarks* to measure μ s necessary for each operation, and comparison of different implementations**

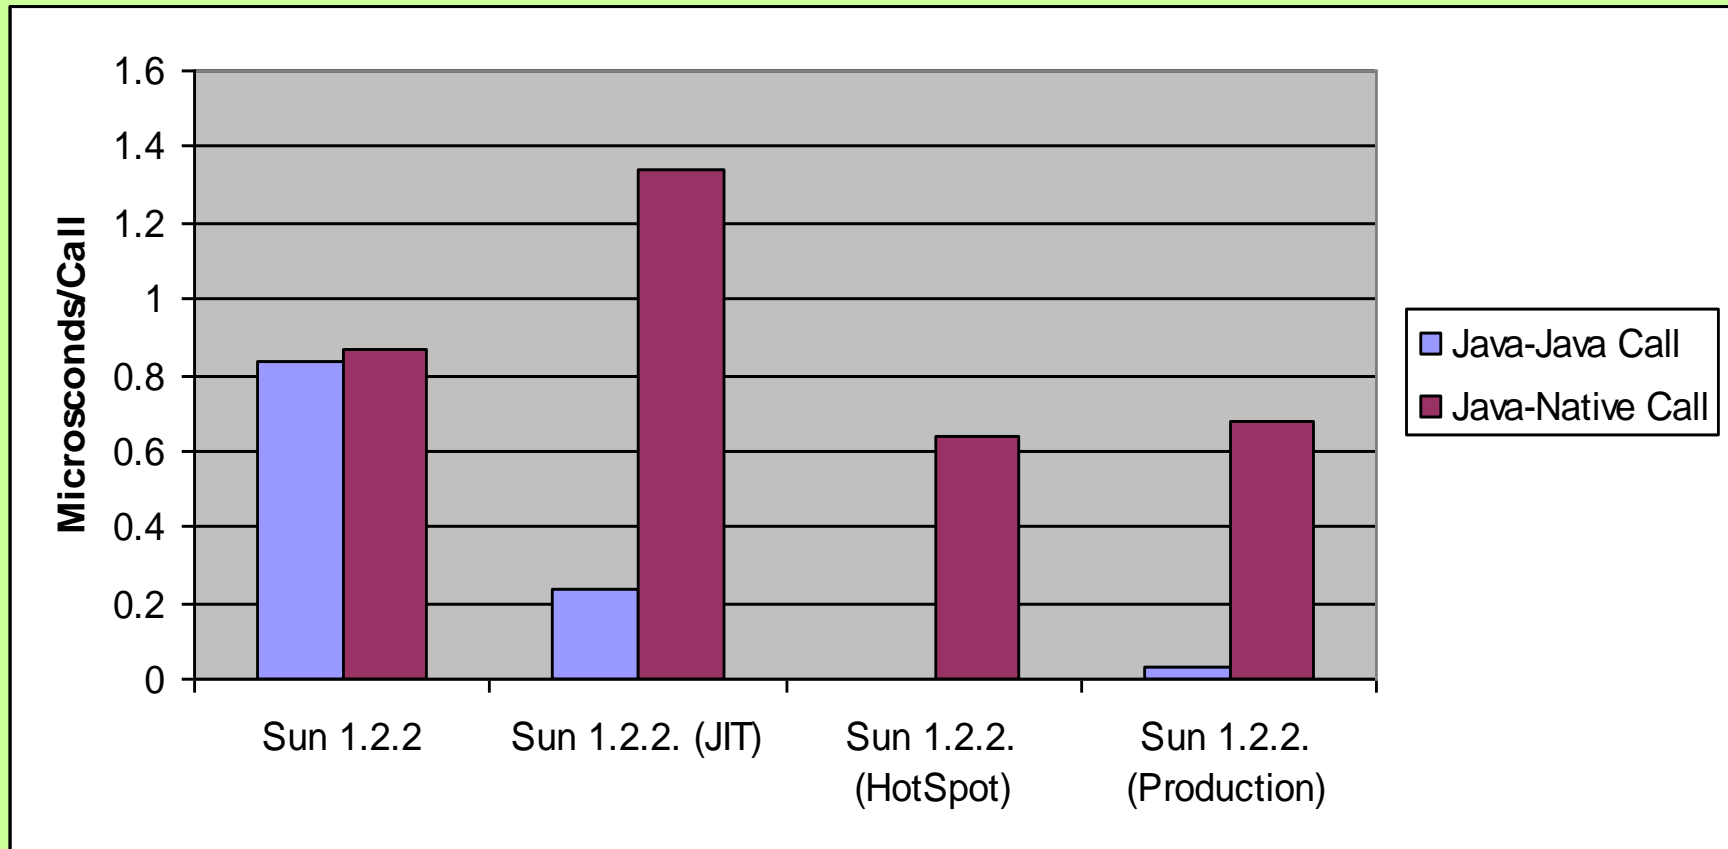
JNI Performance

Native Call Cost: SGI Results




JNI Performance

Native Call Cost: Sun Results



JNI Performance

Native Call Cost: Linux Results



JNI Performance

Native Call Cost: Conclusions

- **Native method overhead is significant**
- **Overhead has decreased from JDK 1.1.x to Java2**
- **JIT compilation can make JNI calls...**
 - **significantly faster than under the interpreter (SGI Java2, IBM 1.1.8, HotSpot)**
 - **significantly slower (Blackdown Java2)**

More Benchmark references

For you observations

JNI Benchmarking

Java Timer Class from UCSD Benchmark

```
public class Timer
{
    /* various field and method definitions ... */
    public void mark() {
        base_time = System.currentTimeMillis();
    }
    public void record() {
        elapsed_time += (System.currentTimeMillis() -
                        base_time);
    }
    public float elapsed() {
        return ((float) elapsed_time) / 1000;
    }
}
```

JNI Benchmarking

Use of Timer Class

```
t.mark() ;  
longRunningTest() ;  
t.record() ;  
System.err.println(name + " " +  
                    t.elapsed());
```

```
for (int i=0; i<iterations; i++) {  
    t.mark() ;  
    longRunningTest() ;  
    t.record() ;  
}  
System.err.println(name + " " +  
                    (t.elapsed() / iterations));
```

If the test duration is much larger than the granularity of the timer...

JNI Benchmarking

Use of Timer Class

```
controlTimer.mark();
for (int i=0; i<iterations; i++) {
    // unremovable code ...
}
controlTimer.record();

t.mark();
for (int i=0; i<iterations; i++) {
    // unremovable code ...
    veryQuickTest();
}
t.record();

System.err.println(name + " " +
    ((t.elapsed() - controlTimer.elapsed())
    * (1000000.0 / iterations)));
```

Typical durations of our JNI tests (in μs) are much smaller than the granularity of `currentTimeMillis()` (ms), so we measure many executions inside a loop and then account for the loop overhead


JNI Benchmarking

Platforms Tested

- **SGI 2000 (300 MHz R10000), IRIX 6.5.7**
 - **JDK 1.1.8, Java2 1.2.2**
- **Dell (350 MHz Pentium II), Red Hat Linux**
 - **IBM JDK 1.1.8, Blackdown JDK 1.1.7, Blackdown Java2-prev2, Sun Java2 1.3 Beta (including HotSpot Client VM)**
- **Sun UltraSPARC (143 MHz)**
 - **Sun JDK 1.2.2, HotSpot Server 1.0**


JNI Performance

Native Call Cost: SGI Results




JNI Performance

Native Call Cost: Sun Results



JNI Performance

Native Call Cost: Linux Results




JNI Performance

Native Call Cost: Conclusions

- **Native method overhead is significant**
- **Overhead has decreased from JDK 1.1.x to Java2**
- **JIT compilation can make JNI calls...**
 - **significantly faster than under the interpreter (SGI Java2, IBM 1.1.8, HotSpot)**
 - **significantly slower (Blackdown Java2)**

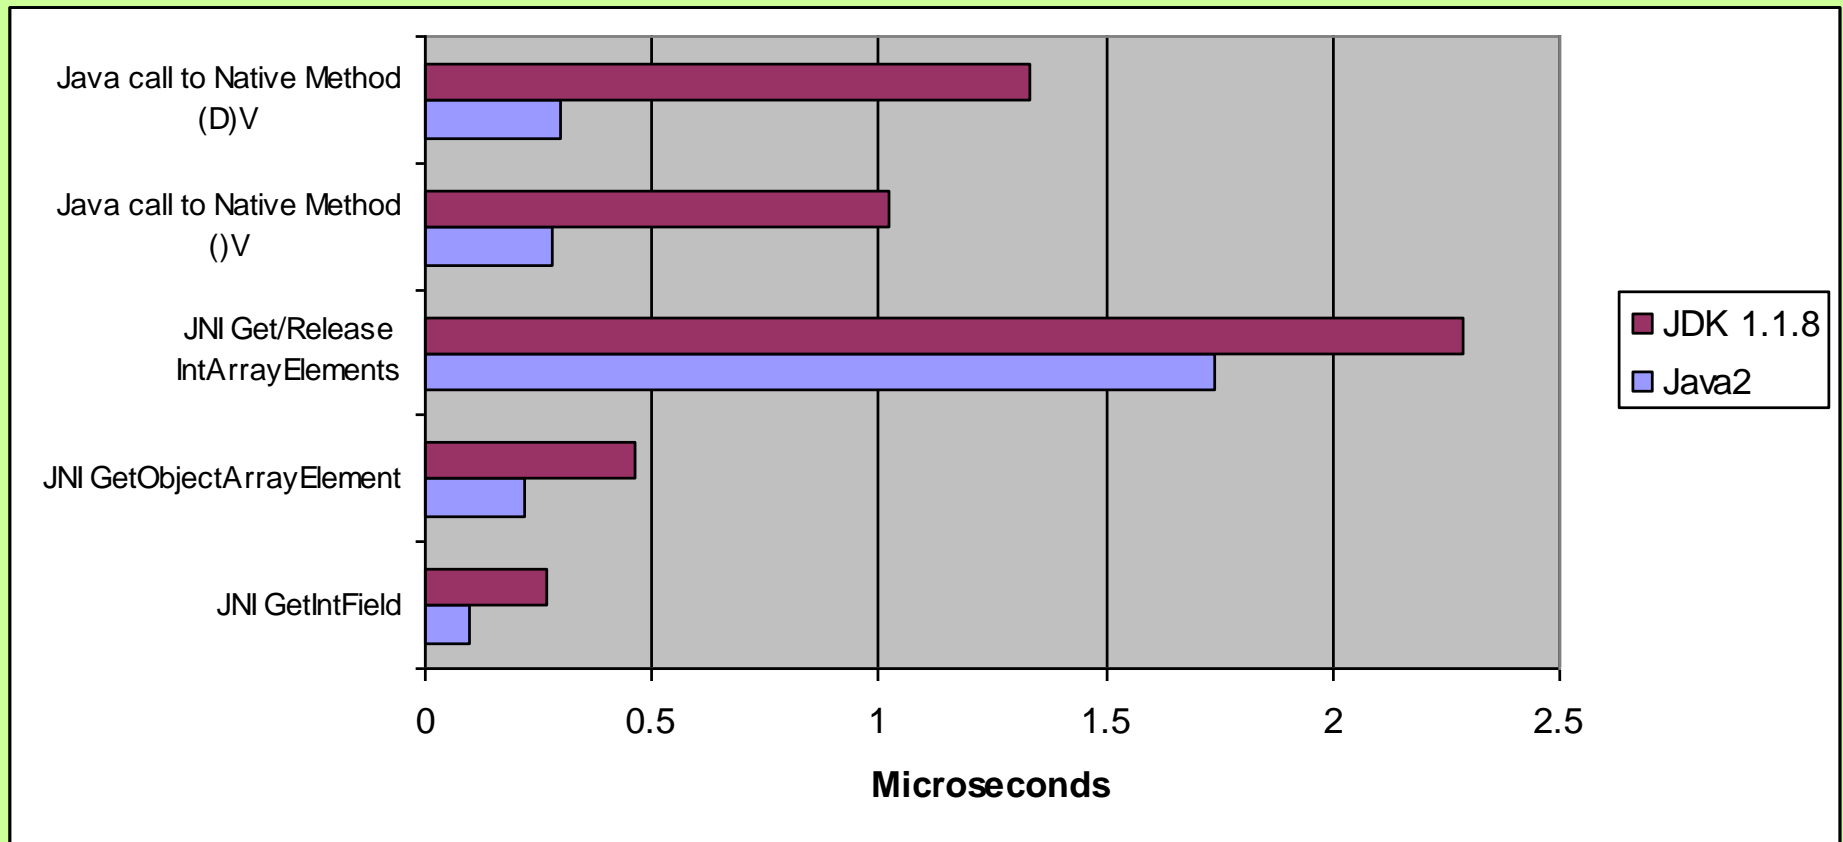
JNI Performance

JNI Performance Landscape



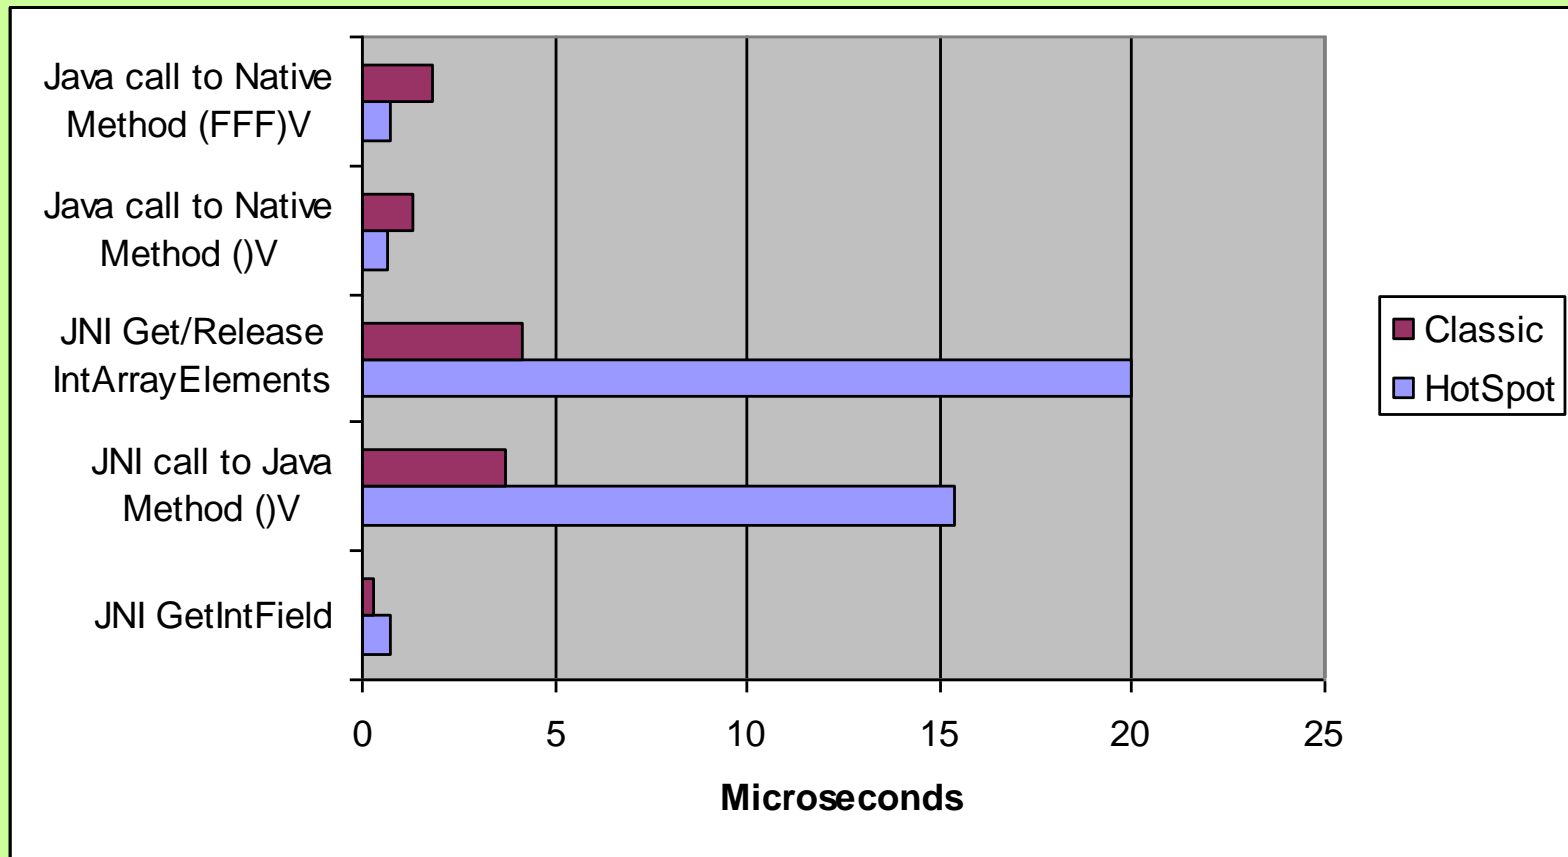
JNI Performance

JDK 1.1.8 v. Java2 (SGI)



JNI Performance

Classic v. HotSpot (Sun Java2)



JNI Performance

JNI Performance: Conclusions

- **Native calls and scalar parameter passing can be relatively efficient with JIT support**
- **Native access to Java arrays can be costly, even if the JVM uses pinning instead of copying**
- **JNI Global References are costly**
- **Calls back into Java are costly**
- **JNI implementations are mostly improving**

JNI Performance

Coding Recommendations

- **Choose the right granularity**
- **Partition application to use native code well**
 - **Good: computation**
 - **Bad: accesses to Java fields and methods**
 - **Necessary: low-level system access**
- **Caching: array data, field and method ID's**

Conclusions

- **Overhead of native method calls and JNI functionality is decreasing, but still deserves careful attention from developers**
- **Our IRIX/MIPS Fast JNI implementation has improved the performance of real-world applications such as Prowess and the Java OpenGL bindings**